

# Markov Chain Monte Carlo and Stan

## Lecture 3

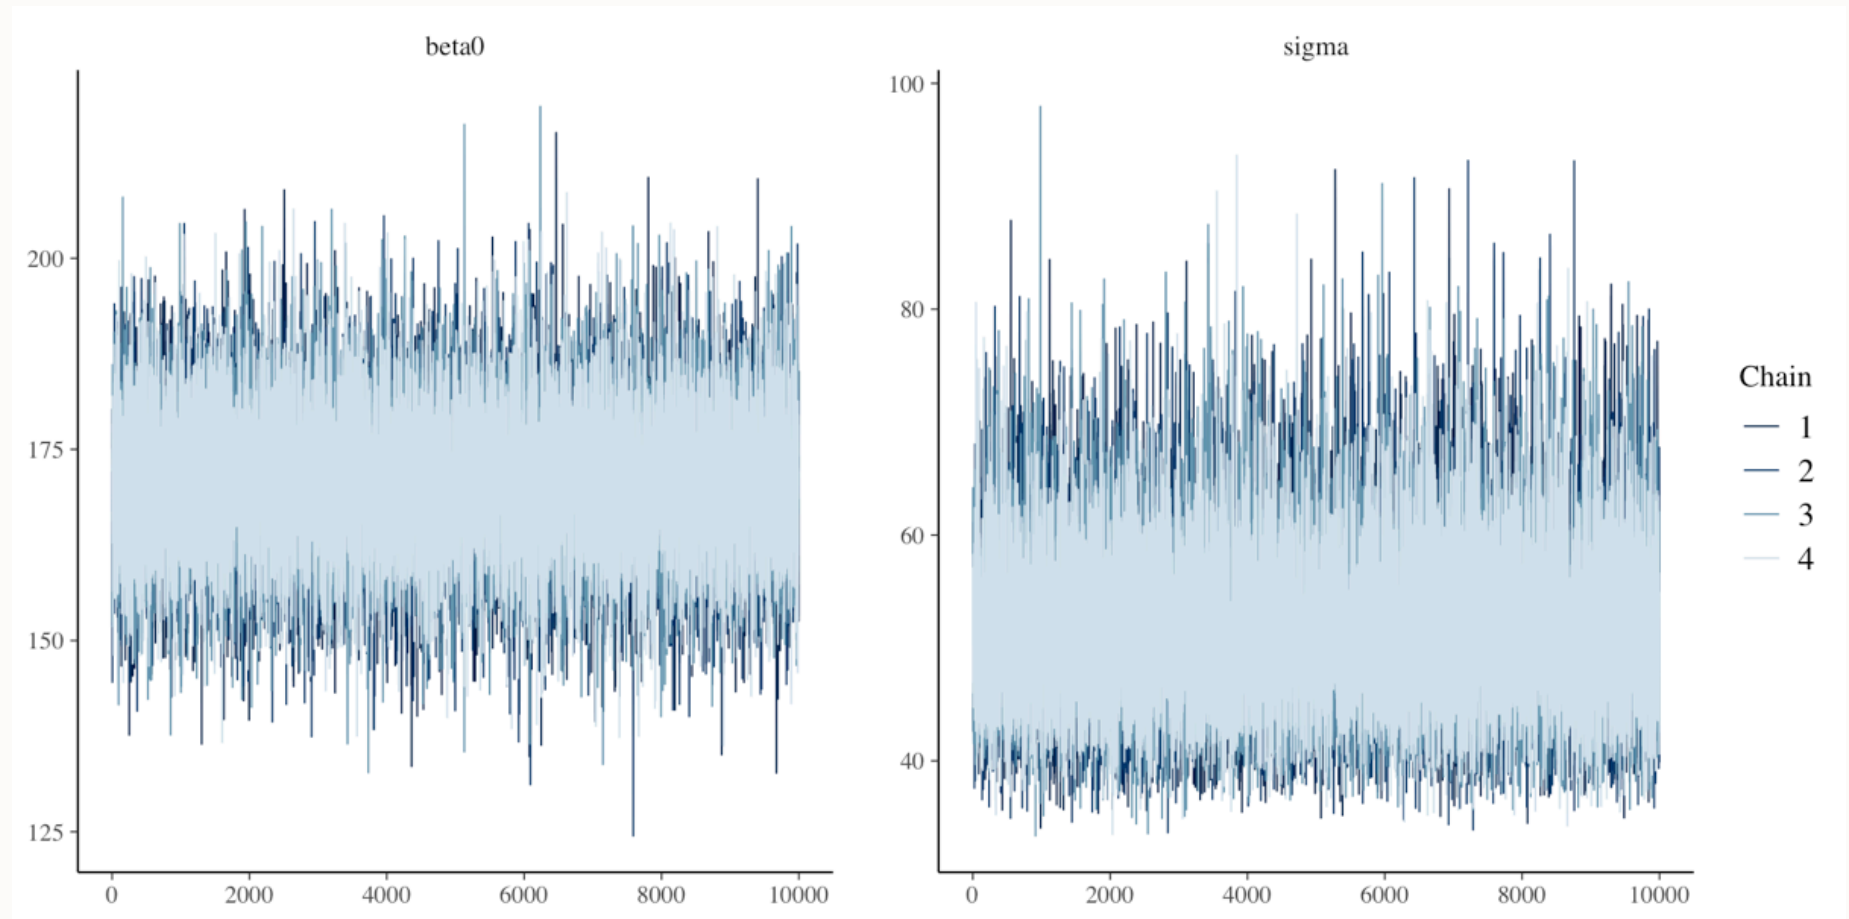
<https://jonathantemplin.com/bayesian-psychometric-modeling-fall-2024/>

1

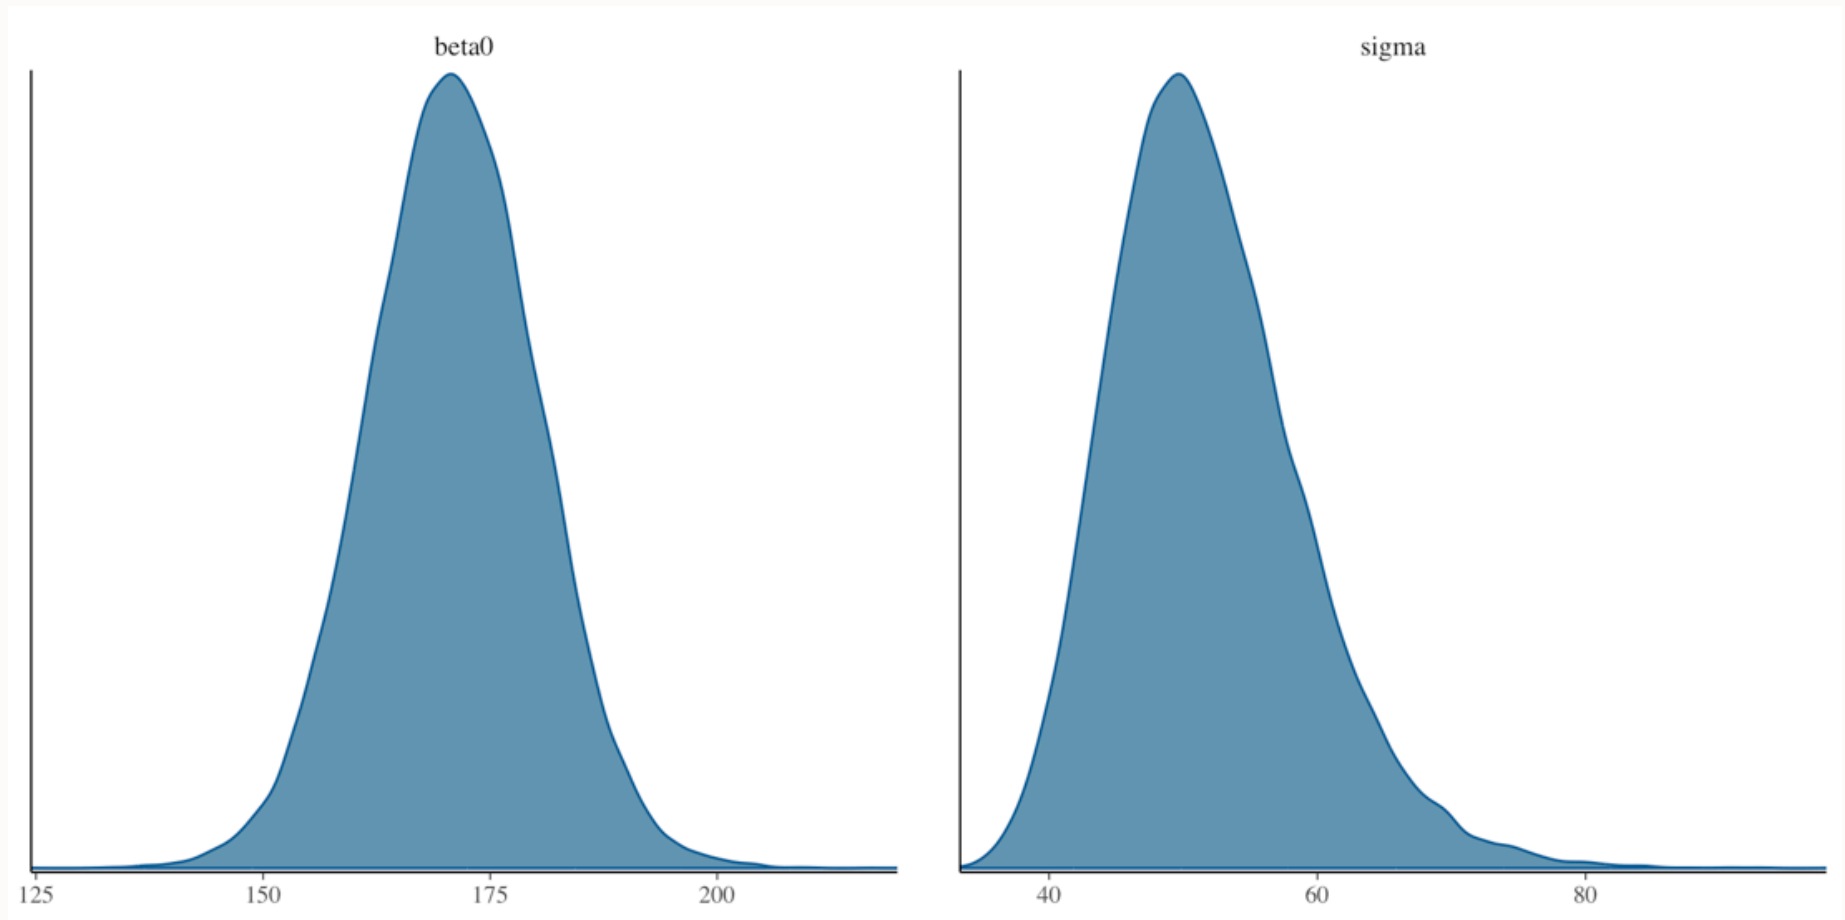
# Today's Lecture Objectives

1. An Introduction to MCMC
2. An Introduction to Stan
3. Both with Linear Models

# The Markov Chain Timeseries



# The Posterior Distribution



# Markov Chain Monte Carlo Estimation

Bayesian analysis is all about estimating the posterior distribution

- Up until now, we have worked with posterior distributions that fairly well-known
  - Beta-Binomial had a Beta distribution
  - In general, likelihood distributions from the exponential family have conjugate priors
    - Conjugate prior: the family of the prior is equivalent to the family of the posterior
- Most of the time, however, posterior distributions are not easily obtainable
  - No longer able to use properties of the distribution to estimate parameters

# Markov Chain Monte Carlo Estimation

- It is possible to use an optimization algorithm (e.g., Newton-Raphson or Expectation-Maximization) to find maximum value of posterior distribution
  - But, such algorithms may take a very long time for high-dimensional problems
- Variational Bayesian methods also attempt to approximate the center of the posterior distribution (and are faster than optimization algorithms for large models)
- Instead: MCMC “sketches” the posterior by sampling from it — then use that sketch to make inferences
  - Sampling is done via MCMC

# Markov Chain Monte Carlo Estimation

- MCMC algorithms iteratively sample from the posterior distribution
  - For fairly simplistic models, each iteration has independent samples
  - Most models have some layers of dependency included
    - Can slow down sampling from the posterior
- There are numerous variations of MCMC algorithms
  - Most of these specific algorithms use one of two types of sampling:
    1. Direct sampling from the posterior distribution (i.e. Gibbs sampling)
      - Often used when conjugate priors are specified
    2. Indirect (rejection-based) sampling from the posterior distribution (e.g., Metropolis-Hastings, Hamiltonian Monte Carlo)

# MCMC Algorithms

- Efficiency is the main reason for so many algorithms
  - Efficiency in this context: How quickly the algorithm converges and provides adequate coverage (“sketching”) of the posterior distribution
  - No one algorithm is uniformly most efficient for all models (here model = likelihood prior)
- The good news is that many software packages (stan, JAGS, MPlus, especially) don’t make you choose which specific algorithm to use
- The bad news is that sometimes your model may take a large amount of time to reach convergence (think days or weeks)
- You can also code your own custom algorithm to make things run smoother



# Commonalities Across MCMC Algorithms

- Despite having fairly broad differences regarding how algorithms sample from the posterior distribution, there are quite a few things that are similar across algorithms:
  1. A period of the Markov chain where sampling is not directly from the posterior
    - The burnin period (sometimes coupled with other tuning periods and called warmup)
  2. Methods used to assess convergence of the chain to the posterior distribution
    - Often involving the need to use multiple chains with independent and differing starting values
  3. Summaries of the posterior distribution

# Commonalities Across MCMC Algorithms

- Further, rejection-based sampling algorithms often need a tuning period to make the sampling more efficient
  - The tuning period comes before the algorithm begins its burnin period

# MCMC Demonstration

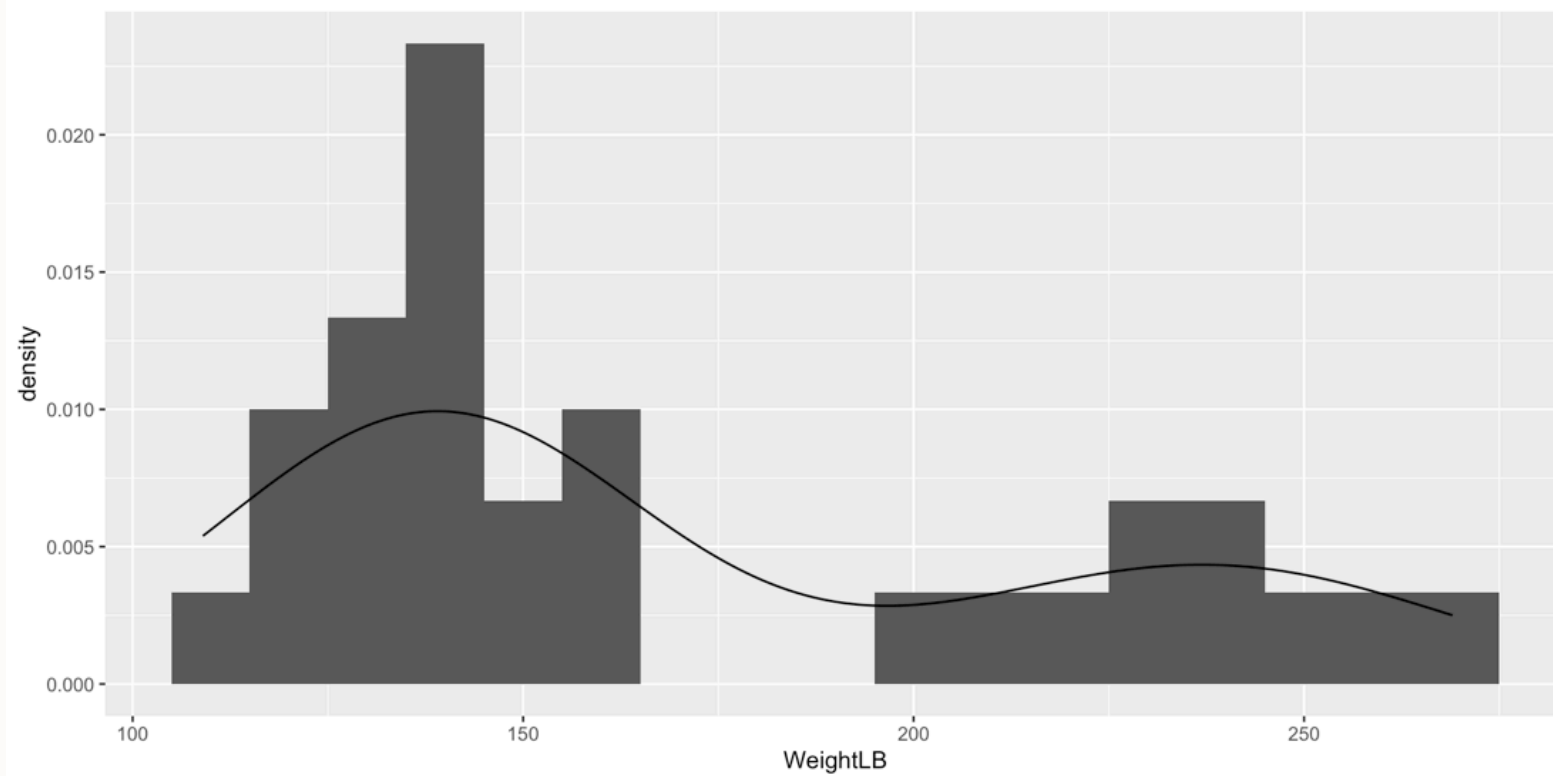
- To demonstrate each type of algorithm, we will use a model for a normal distribution
  - We will investigate each, briefly
  - We will then switch over to stan to show the syntax and let stan work
  - We will conclude by talking about assessing convergence and how to report parameter estimates.

# Example Data: Post-Diet Weights

[Example Data Link](#)

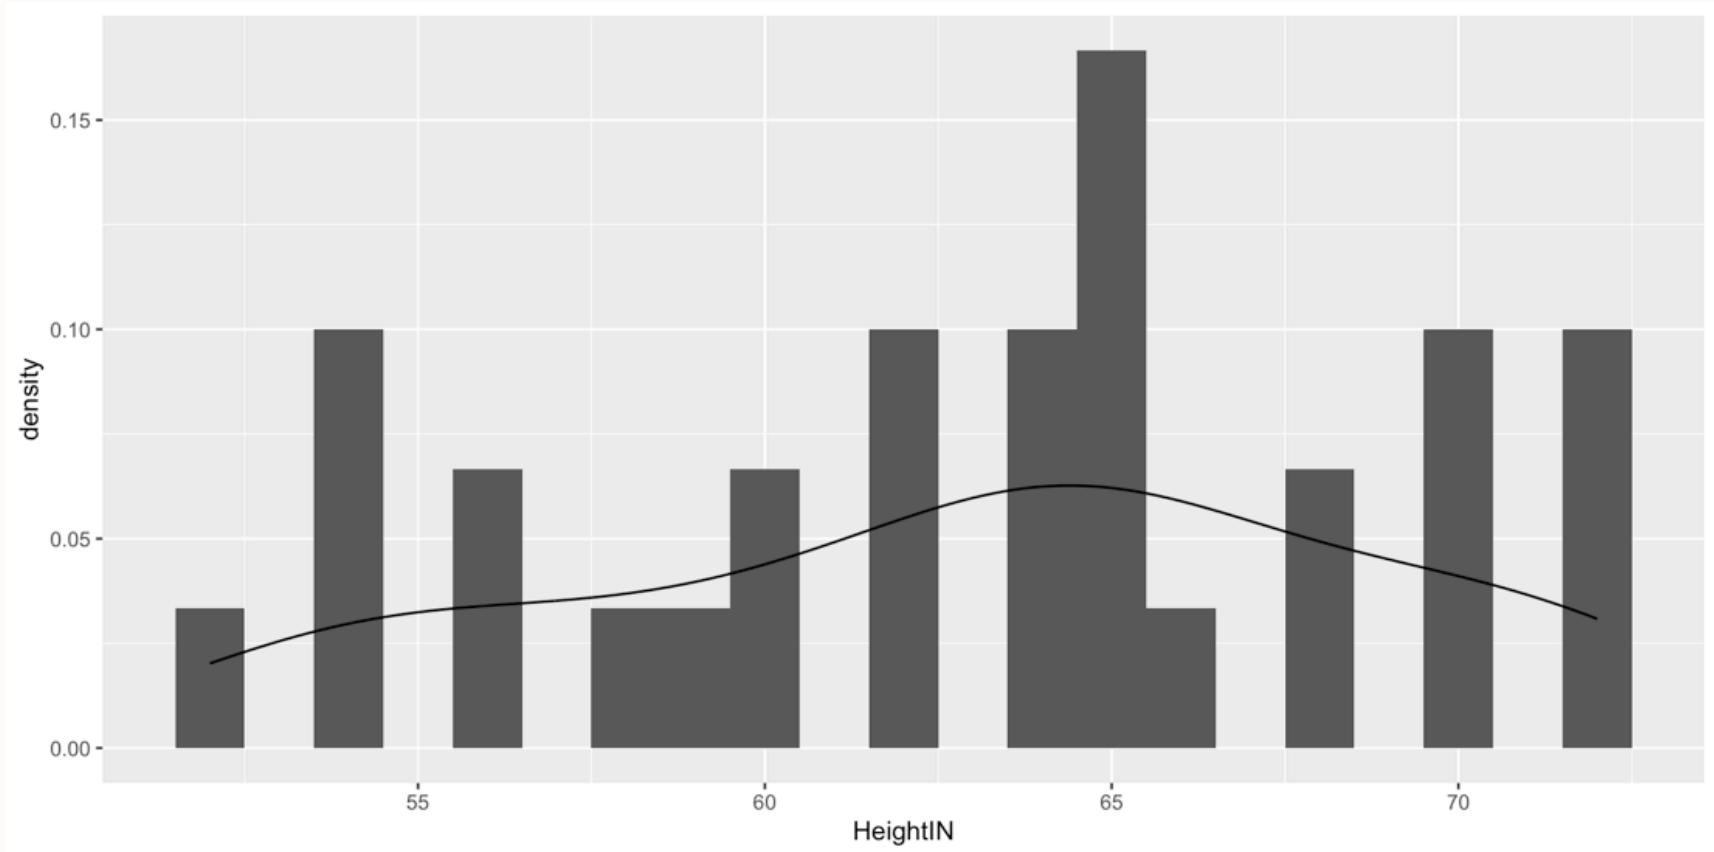
# Visualizing Data: WeightLB Variable

```
1 DietData = read.csv(file = "DietData.csv")
2
3 ggplot(data = DietData, aes(x = WeightLB)) +
4   geom_histogram(aes(y = ..density..), position = "identity", binwidth = 10) +
5   geom_density(alpha=.2)
```



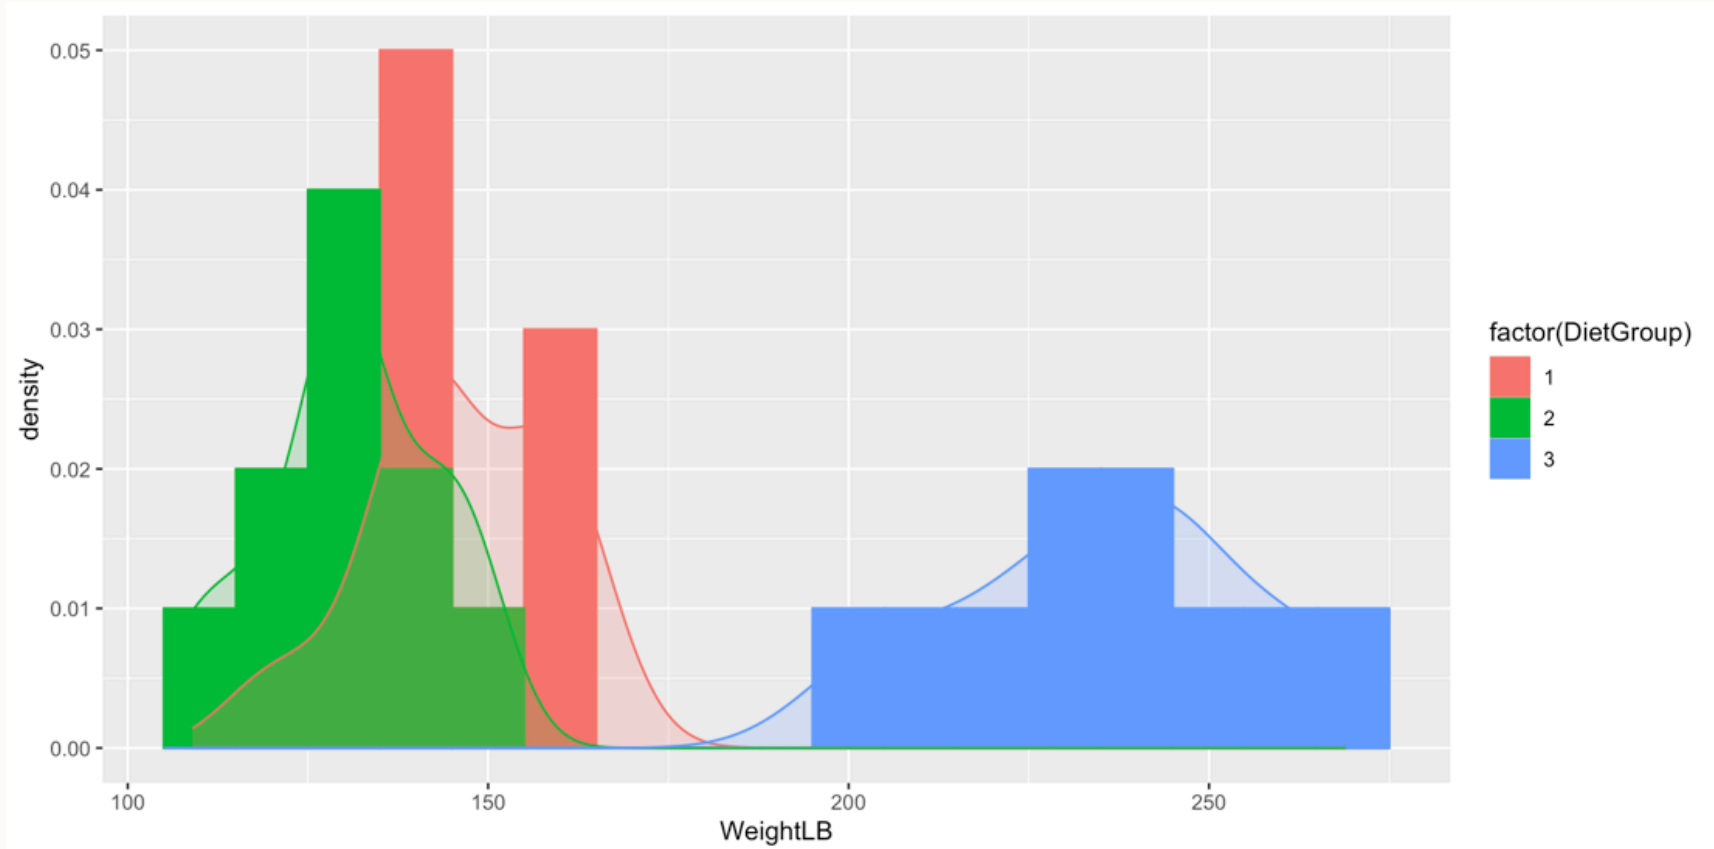
# Visualizing Data: HeightIN Variable

```
1 ggplot(data = DietData, aes(x = HeightIN)) +  
2   geom_histogram(aes(y = ..density..), position = "identity", binwidth = 1) +  
3   geom_density(alpha=.2)
```



# Visualizing Data: WeightLB by Group

```
1 ggplot(data = DietData, aes(x = WeightLB, color = factor(DietGroup), fill = factor(DietGroup))) +  
2   geom_histogram(aes(y = ..density..), position = "identity", binwidth = 10) +  
3   geom_density(alpha=.2)
```



# Visualizing Data: Weight by Height by Group

```
1 ggplot(data = DietData, aes(x = HeightIN, y = WeightLB, shape = factor(DietGroup), color = factor(Die
2   geom_smooth(method = "lm", se = FALSE) + geom_point())
```



# Class Discussion: What Do We Do?

Now, your turn to answer questions:

1. What type of analysis seems most appropriate for these data?
2. Is the dependent variable (`WeightLB`) is appropriate as-is for such analysis or does it need transformed?

# Linear Model with Least Squares

Let's play with models for data...

```

1 # center predictors for reasonable numbers
2 DietData$HeightIN60 = DietData$HeightIN-60
3
4 # full analysis model suggested by data:
5 FullModel = lm(formula = WeightLB ~ 1, data = DietData)
6
7 # examining assumptions and leverage of fit
8 # plot(FullModel)
9
10 # looking at ANOVA table
11 # anova(FullModel)
12
13 # looking at parameter summary
14 summary(FullModel)

```

Call:

```
lm(formula = WeightLB ~ 1, data = DietData)
```

Residuals:

Min	1Q	Median	3Q	Max
-62.00	-36.75	-24.00	49.00	98.00

Coefficients:

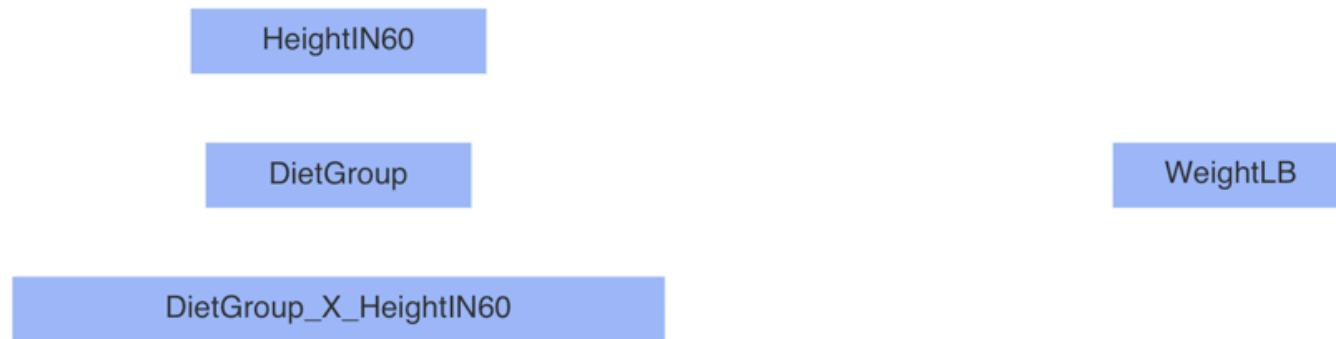
	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	171.000	9.041	18.91	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 49.52 on 29 degrees of freedom

# Path Diagram of Our Model



# Steps in an MCMC Analysis

1. Specify model
2. Specify prior distributions for all model parameters
3. Build model syntax as needed
4. Run Markov chains (specify warmup/burnin and sampling period lengths)
5. Evaluate chain convergence
6. Interpret/report results

# Specify Model

- To begin, let's start with an empty model and build up from there
- Let's examine the linear model we seek to estimate:

Where:


Questions:

- What are the *variables* in this analysis?
- What are the parameters in this analysis?

# Introduction to Stan

- Stan is an MCMC estimation program
  - Most recent; has many convenient features
  - Actually does several methods of estimation (ML, Variational Bayes)
- You create a model using Stan's syntax
  - Stan translates your model to a custom-built C++ syntax
  - Stan then compiles your model into its own executable program
- You then run the program to estimate your model
  - If you use R, the interface can be seamless

# Stan and RStudio

- Stan has its own syntax which can be built in stand-alone text files 
  - Rstudio will let you create one of these files in the new file menu
  - Rstudio also has syntax highlighting in Stan files
    - This is very helpful to learn the syntax
- Stan syntax can also be built from R character strings
  - Which is helpful when running more than one model per analysis

# Stan Syntax

```
1 data {  
2   int<lower=0> N;  
3   vector[N] y;  
4 }  
5  
6 parameters {  
7   real beta0;  
8   real<lower=0> sigma;  
9 }  
10  
11 model {  
12   beta0 ~ normal(0, 1000); // prior for beta0  
13   sigma ~ uniform(0, 100000); // prior for sigma  
14   y ~ normal(beta0, sigma); // model for observed data  
15 }
```

- Above is the syntax for our model
  - Each line ends with a semi colon
  - Comments are put in with //



# Stan Syntax: R Character Object

```
1 stanModel = "  
2 data {  
3   int<lower=0> N;  
4   vector[N] y;  
5 }  
6  
7 parameters {  
8   real beta0;  
9   real<lower=0> sigma;  
10 }  
11  
12 model {  
13   beta0 ~ normal(0, 1000); // prior for beta0  
14   sigma ~ uniform(0, 100000); // prior for sigma  
15   y ~ normal(beta0, sigma); // model for observed data  
16 }  
17 "
```

- Three blocks of syntax needed
  - Data: What Stan expects you will send to it for the analysis (using R lists)
  - Parameters: Where you specify what the parameters of the model are
  - Model: Where you specify the distributions of the priors and data

# Stan Data and Parameter Declaration

Like many compiled languages, Stan expects you to declare what type of data/parameters you are defining:

- `int`: Integer values (no decimals)
- `real`: Floating point numbers
- `vector`: A one-dimensional set of real valued numbers

Sometimes, additional definitions are provided giving the range of the variable (or restricting the set of starting values):

- `real<lower=0> sigma;`

See: <https://mc-stan.org/docs/reference-manual/data-types.html> for more information

# Stan Data and Prior Distributions

- In the model section, you define the distributions needed for the model and the priors
  - The left-hand side is either defined in data or parameters
    - `y ~ normal(beta0, sigma); // model for observed data`
    - `sigma ~ uniform(0, 100000); // prior for sigma`
  - The right-hand side is a distribution included in Stan
    - You can also define your own distributions

See: <https://mc-stan.org/docs/functions-reference/index.html> for more information

# From Stan Syntax to Compilation

```
1 # compile model -- this method is for stand-alone stan files (uses cmdstanr)
2 model00.fromFile = cmdstan_model(stan_file = "model00.stan")
3
4 # or this method using the string text in R
5 model00.fromString = cmdstan_model(stan_file = write_stan_file(stanModel))
```

- Once you have your syntax, next you need to have Stan translate it into C++ and compile an executable
- This is where `cmdstanr` and `rstan` differ
  - `cmdstanr` wants you to compile first, then run the Markov chain
  - `rstan` conducts compilation (if needed) then runs the Markov chain

# Building Data for Stan

```
1 # build R list containing data for Stan: Must be named what "data" are listed in analysis
2 stanData = list(
3   N = nrow(DietData),
4   y = DietData$WeightLB
5 )
6
7 # snippet of Stan syntax:
8 stanSyntaxSnippet = "
9 data {
10   int<lower=0> N;
11   vector[N] y;
12 }
13 "
```

- Stan needs the data you declared in your syntax to be able to run
- Within R, we can pass this data to Stan via a list object
- The entries in the list should correspond to the data portion of the Stan syntax
  - In the above syntax, we told Stan to expect a single integer named **N** and a vector named **y**
- The R list object is the same for **cmdstanr** and **rstan**

# Running Markov Chains in cmdstanr

```
1 # run MCMC chain (sample from posterior)
2 model00.samples = model00.fromFile$sample(
3   data = stanData,
4   seed = 1,
5   chains = 4,
6   parallel_chains = 4,
7   iter_warmup = 10000,
8   iter_sampling = 10000
9 )
```

# Running Markov Chains in `rstan`

```
1 rstan_options(auto_write = TRUE)
2 options(mc.cores = parallel::detectCores())
3
4 # example MCMC analysis in rstan
5 model00.rstan = stan(
6   model_code = stanModel,
7   model_name = "Empty model",
8   data = stanData,
9   warmup = 10000,
10  iter = 20000,
11  chains = 4,
12  verbose = TRUE
13 )
```

- Rstan takes the model syntax directly, then compiles and runs the chains
- The first two lines of syntax enable running one chain per thread (parallel processing)
  - As chains are independent, running them simultaneously (parallel) shortens wait time considerably
- The `verbose` option is helpful for detecting when things break
- The same R list supplies the data to Stan

# MCMC Process

- The MCMC algorithm runs as a series of discrete iterations
  - Within each iteration, each parameter of a model has an opportunity to change its value
- For each parameter, a new parameter is sampled at random from the current belief of posterior distribution
  - The specifics of the sampling process differ by algorithm type (we'll have a lecture on this later)



# MCMC Process

- In Stan (Hamiltonian Monte Carlo), for a given iteration, a proposed parameter is generated
  - The posterior likelihood “values” (more than just density; includes likelihood of proposal) are calculated for the current and proposed values of the parameter
  - The proposed values are accepted based on the draw of a uniform number compared to a transition probability
- If all models are specified correctly, then regardless of starting location, each chain will converge to the posterior if run long enough
  - But, the chains must be checked for convergence when the algorithm stops

# Example of Bad Convergence

# Examining Chain Convergence

- Once Stan stops, the next step is to determine if the chains converged to their posterior distribution
  - Called convergence diagnosis
- Many methods have been developed for diagnosing if Markov chains have converged
  - Two most common: visual inspection and Gelman-Rubin Potential Scale Reduction Factor (PSRF; [quick reference](#))

# Examining Chain Convergence

- Visual inspection
  - Want no trends in timeseries — should look like a caterpillar
  - Shape of posterior density should be mostly smooth
- Gelman-Rubin PSRF (denoted with  $\hat{R}$ )
  - For analyses with multiple chains
  - Ratio of between-chain variance to within-chain variance
  - Should be near 1 (maximum somewhere under 1.1)

# Setting MCMC Options

- As convergence is assessed using multiple chains, more than one should be run
  - Between-chain variance estimates improve with the number of chains, so I typically use four
  - Others have two; more than one should work
- Warmup/burnin period should be long enough to ensure chains move to center of posterior distribution
  - Difficult to determine ahead of time
  - More complex models need more warmup/burnin to converge
- Sampling iterations should be long enough to thoroughly sample posterior distribution
  - Difficulty to determine ahead of time
  - Need smooth densities across bulk of posterior
- Often, multiple analyses (with different settings) is what is needed

# The Markov Chain Timeseries

# The Posterior Distribution

# Assessing Our Chains

```
1 model00.samples$summary()

# A tibble: 3 × 10
  variable    mean median    sd   mad    q5   q95  rhat ess_bulk ess_tail
  <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
1 lp__      -129.  -128.  1.02 0.735 -131. -128.  1.00   17577.   22523.
2 beta0     171.   171.  9.44 9.31  155.  186.  1.00   29639.   25175.
3 sigma     51.8   51.0  7.20 6.84   41.5  64.8  1.00   29147.   24620.
```

- The summary function reports the PSRF (rhat)
- Here we look at our two parameters: and
- Both have , so both would be considered converged
- `lp__` is posterior log likelihood—does not necessarily need examined
- `ess_` columns show effect sample size for chain (factoring in autocorrelation between correlations)
  - More is better



# Results Interpretation

- At long last, with a set of converged Markov chains, we can now interpret the results
  - Here, we disregard which chain samples came from and pool all sampled values to use for results
- We use summaries of posterior distributions when describing model parameters
  - Typical summary: the posterior mean
    - The mean of the sampled values in the chain
  - Called EAP (Expected a Posteriori) estimates
  - Less common: posterior median

# Results Interpretation

- Important point:
  - Posterior means are different than what characterizes the ML estimates
    - Analogous to ML estimates would be the mode of the posterior distribution
  - Especially important if looking at non-symmetric posterior distributions
    - Look at posterior for variances

# Results Interpretation

- To summarize the uncertainty in parameters, the posterior standard deviation is used
  - The standard deviation of the sampled values in the chain
  - This is the analogous to the standard error from ML
- Bayesian credible intervals are formed by taking quantiles of the posterior distribution
  - Analogous to confidence intervals
  - Interpretation slightly different – the probability the parameter lies within the interval
  - 95% credible interval notes that parameter is within interval with 95% confidence
- Additionally, highest density posterior intervals can be formed
  - The narrowest range for an interval (for unimodal posterior distributions)

# Our Results

```
1 model00.samples$summary()
```

```
# A tibble: 3 × 10
```

	variable	mean	median	sd	mad	q5	q95	rhat	ess_bulk	ess_tail
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	lp__	-129.	-128.	1.02	0.735	-131.	-128.	1.00	17577.	22523.
2	beta0	171.	171.	9.44	9.31	155.	186.	1.00	29639.	25175.
3	sigma	51.8	51.0	7.20	6.84	41.5	64.8	1.00	29147.	24620.

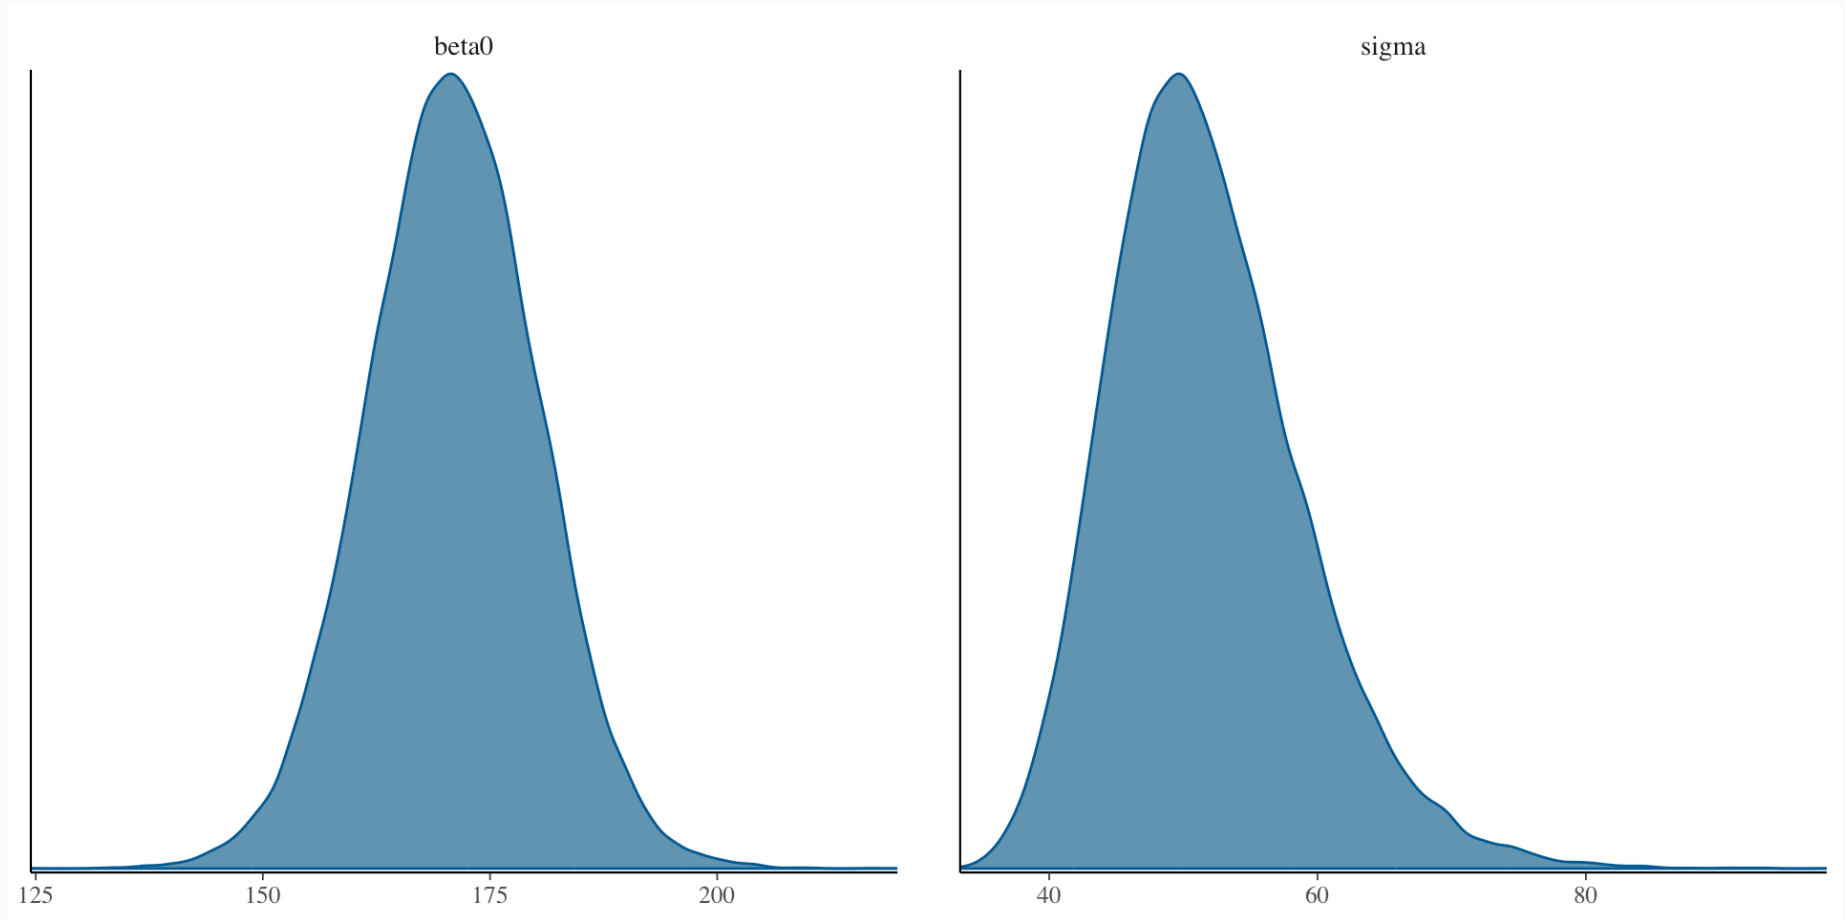
```
1 hdi(model00.samples$draws("beta0"), credMass = .9)
```

```
lower upper
155.015 185.762
attr(,"credMass")
[1] 0.9
```

```
1 hdi(model00.samples$draws("sigma"), credMass = .9)
```

```
lower upper
40.2305 63.0083
attr(,"credMass")
[1] 0.9
```

# The Posterior Distribution



# Wrapping Up

- This lecture covered the basics of MCMC estimation with Stan
- Next we will use an example to show a full analysis of the data problem we started with today
- The details today are the same for all MCMC analyses, regardless of which algorithm is used