

# Efficient Stan Code and Generated Quantities

## Lecture 3b

<https://jonathantemplin.com/bayesian-psychometric-modeling-fall-2024/>

1

# Today's Lecture Objectives

1. Making Stan Syntax Shorter
2. Computing Functions of Model Parameters

# Making Stan Code Shorter

The Stan syntax from our previous model was lengthy:

- A declared variable for each parameter
- The linear combination of coefficients multiplying predictors

Stan has built-in features to shorten syntax:

- Matrices/Vectors
- Matrix products
- Multivariate distributions (initially for prior distributions)

# Linear Models without Matrices

The linear model from our example was:

$$\text{WeightLB}_p = \beta_0 + \beta_1 \text{HeightIN}_p + \beta_2 \text{Group2}_p + \beta_3 \text{Group3}_p + \beta_4 \text{HeightIN}_p \text{Group2}_p - \beta_5 \text{HeightIN}_p \text{Group3}_p + e_p$$

with:

- $\text{Group2}_p$  the binary indicator of person  $p$  being in group 2
- $\text{Group3}_p$  the binary indicator of person  $p$  being in group 3
- $e_p \sim N(0, \sigma_e)$

# Path Diagram of Model



# Linear Models with Matrices

Model (predictor) matrix:

$$\mathbf{X} = \begin{bmatrix} 1 & -4 & 0 & 0 & 0 & 0 \\ & & \vdots & & & \\ 1 & 12 & 0 & 1 & 0 & 12 \end{bmatrix}$$

Coefficients vector:

$$\boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \\ \beta_4 \\ \beta_5 \end{bmatrix}$$

```
1 head(model.matrix(FullModelFormula, data = DietData))
(Intercept) Height60IN factor(DietGroup)2 factor(DietGroup)3
1          1         -4                0                0
2          1          0                0                0
3          1          4                0                0
4          1          8                0                0
5          1         12                0                0
6          1         -6                0                0
Height60IN:factor(DietGroup)2 Height60IN:factor(DietGroup)3
1                0                0
2                0                0
3                0                0
4                0                0
5                0                0
```



# Linear Models with Matrices

Using matrices, we can rewrite our regression equation from

$$\text{WeightLB}_p = \beta_0 + \beta_1 \text{HeightIN}_p + \beta_2 \text{Group2}_p + \beta_3 \text{Group3}_p + \beta_4 \text{HeightIN}_p \text{Group2}_p - \beta_5 \text{HeightIN}_p \text{Group3}_p + e_p$$

To:

$$\mathbf{WeightLB} = \mathbf{X}\boldsymbol{\beta} + \mathbf{e}$$

Where:



# Example: Predicted Values

```

1 P=6
2 beta = matrix(data = runif(n = 6, min = 0, max = 10), nrow = P, ncol = 1)
3 X = model.matrix(FullModelFormula, data=DietData)
4 X %*% beta # R uses %*% for matrix products

```

```

      [,1]
1    3.5041870
2    4.2407897
3    4.9773925
4    5.7139952
5    6.4505980
6    3.1358856
7    4.6090911
8    5.1615432
9    5.1615432
10   6.0822966
11  -1.5318296
12   5.4390386
13  12.4099067
14  19.3807748
15  26.3516430
16  -5.0172636
17   8.9244726
18  14.1526237
19  14.1526237
20  22.8662089
21  25.4120000

```

# Syntax Changes: Data Section

## Old Syntax Without Matrices

```
1 data {  
2   int<lower=0> N;  
3   vector[N] weightLB;  
4   vector[N] height60IN;  
5   vector[N] group2;  
6   vector[N] group3;  
7   vector[N] heightXgroup2;  
8   vector[N] heightXgroup3;  
9 }
```

## New Syntax With Matrices

```
1 data {  
2   int<lower=0> N;           // number of observations  
3   int<lower=0> P;           // number of predictors (plus column for intercept)  
4   matrix[N, P] X;          // model.matrix() from R  
5   vector[N] y;             // outcome  
6  
7   vector[P] meanBeta;       // prior mean vector for coefficients  
8   matrix[P, P] covBeta;    // prior covariance matrix for coefficients  
9  
10  real sigmaRate;           // prior rate parameter for residual standard deviation  
11 }
```

# Syntax Changes: Parameters Section

## Old Syntax Without Matrices

```
1 parameters {  
2   real beta0;  
3   real betaHeight;  
4   real betaGroup2;  
5   real betaGroup3;  
6   real betaHxG2;  
7   real betaHxG3;  
8   real<lower=0> sigma;  
9 }
```

## New Syntax With Matrices

```
1 parameters {  
2   vector[P] beta;           // vector of coefficients for Beta  
3   real<lower=0> sigma;      // residual standard deviation  
4 }
```

# Defining Prior Distributions

Previously, we defined a normal prior distribution for each regression coefficient

- Univariate priors – univariate normal distribution
- Each parameter had a prior that was independent of the other parameters

When combining all parameters into a vector, a natural extension is a multivariate normal distribution

- [https://en.wikipedia.org/wiki/Multivariate\\_normal\\_distribution](https://en.wikipedia.org/wiki/Multivariate_normal_distribution)
- Mean vector (`meanBeta`; size  $P \times 1$ )
  - Put all prior means for these coefficients into a vector from R
- Covariance matrix (`covBeta`; size  $P \times P$ )
  - Put all prior variances (prior  $SD^2$ ) into the diagonal
  - With zeros for off diagonal, the MVN prior is equivalent to the set of independent univariate normal priors

# Syntax Changes: Model Section

## Old Syntax Without Matrices

```
1 model {
2   beta0 ~ normal(0,1);
3   betaHeight ~ normal(0,1);
4   betaGroup2 ~ normal(0,1);
5   betaGroup3 ~ normal(0,1);
6   betaHxG2 ~ normal(0,1);
7   betaHxG3 ~ normal(0,1);
8
9   sigma ~ exponential(.1); // prior for sigma
10  weightLB ~ normal(
11    beta0 + betaHeight * height60IN + betaGroup2 * group2 +
12    betaGroup3 * group3 + betaHxG2 * heightXgroup2 +
13    betaHxG3 * heightXgroup3, sigma);
14 }
```

## New Syntax With Matrices

```
1 model {
2   beta ~ multi_normal(meanBeta, covBeta); // prior for coefficients
3   sigma ~ exponential(sigmaRate); // prior for sigma
4   y ~ normal(X*beta, sigma); // linear model
5 }
```

See: Example Syntax in R File

# Summary of Changes

- With matrices, there is less syntax to write
  - Model is equivalent
- Output, however, is not labeled with respect to parameters
  - May have to label output

```
# A tibble: 8 × 10
  variable    mean  median    sd   mad    q5    q95  rhat ess_bulk ess_tail
  <chr>      <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl>   <dbl>
1 lp__      -78.0   -77.7  2.09  1.93 -82.0 -75.3  1.00   2840.   4327.
2 beta[1]   147.    147.   3.17  3.09 142.  152.  1.00   3044.   4196.
3 beta[2]   -0.349  -0.352 0.485 0.475 -1.15  0.455  1.00   3258.   4432.
4 beta[3]  -24.0   -24.0  4.46  4.40 -31.3 -16.6  1.00   3340.   4801.
5 beta[4]   81.5    81.5  4.22  4.14  74.6  88.5  1.00   3438.   4785.
6 beta[5]    2.45    2.45  0.683 0.680  1.33  3.54  1.00   3579.   4813.
7 beta[6]    3.53    3.53  0.640 0.630  2.48  4.58  1.00   3550.   4266.
8 sigma     8.24    8.10  1.22  1.16  6.51 10.4  1.00   4444.   4860.
```

# Computing Functions of Parameters



# Computing Functions of Parameters

- Often, we need to compute some linear or non-linear function of parameters in a linear model
  - Missing effects (i.e., slope for Diet Group 2)
  - Simple slopes
  - $R^2$
- In non-Bayesian analyses, these are often formed with the point estimates of parameters
- For Bayesian analyses, however, we will seek to build the posterior distribution for any function of the parameters
  - This means applying the function to all posterior samples

# Example: Need Slope for Diet Group 2

Recall our model:

$$\text{WeightLB}_p = \beta_0 + \beta_1 \text{HeightIN}_p + \beta_2 \text{Group2}_p + \beta_3 \text{Group3}_p + \beta_4 \text{HeightIN}_p \text{Group2}_p - \beta_5 \text{HeightIN}_p \text{Group3}_p + e_p$$

Here,  $\beta_1$  is the change in  $\text{WeightLB}_p$  per one-unit change in  $\text{HeightIN}_p$  for a person in Diet Group 1 (i.e.  $\text{Group2}_p$  and  $\text{Group3}_p = 0$ )

Question: What is the slope for Diet Group 2?

- To answer, we need to first form the model when  $\text{Group2}_p = 1$ :

$$\text{WeightLB}_p = \beta_0 + \beta_1 \text{HeightIN}_p + \beta_2 + \beta_4 \text{HeightIN}_p + e_p$$

- Next, we rearrange terms that involve  $\text{HeightIN}_p$ :

$$\text{WeightLB}_p = (\beta_0 + \beta_2) + (\beta_1 + \beta_4) \text{HeightIN}_p + e_p$$

- From here, we can see the slope for Diet Group 2 is  $(\beta_1 + \beta_4)$ 
  - Also, the intercept for Diet Group 2 is  $(\beta_0 + \beta_2)$

# Computing Slope for Diet Group 2

Our task: Create posterior distribution for Diet Group 2

- We must do so for each iteration we've kept from our MCMC chain
- A somewhat tedious way to do this is after using Stan

```
1 model05_Samples$summary()

# A tibble: 8 × 10
  variable    mean median    sd  mad    q5    q95  rhat ess_bulk ess_tail
  <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
1 lp__      -78.0  -77.7  2.09  1.93 -82.0 -75.3  1.00    2840.    4327.
2 beta[1]   147.    147.   3.17  3.09 142.  152.  1.00    3044.    4196.
3 beta[2]   -0.349 -0.352 0.485 0.475 -1.15  0.455  1.00    3258.    4432.
4 beta[3]  -24.0  -24.0  4.46  4.40 -31.3 -16.6  1.00    3340.    4801.
5 beta[4]   81.5   81.5  4.22  4.14  74.6  88.5  1.00    3438.    4785.
6 beta[5]    2.45   2.45  0.683 0.680  1.33  3.54  1.00    3579.    4813.
7 beta[6]    3.53   3.53  0.640 0.630  2.48  4.58  1.00    3550.    4266.
8 sigma     8.24   8.10  1.22  1.16  6.51 10.4  1.00    4444.    4860.

1 slopeG2 = model05_Samples$draws("beta[2]") + model05_Samples$draws("beta[5]")
2 summary(slopeG2)

# A tibble: 1 × 10
  variable    mean median    sd  mad    q5    q95  rhat ess_bulk ess_tail
  <chr>      <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl>
1 beta[2]    2.10   2.10 0.481 0.463  1.31  2.88  1.00    7569.    6251.
```

# Computing the Slope Within Stan

Stan can compute these values for us—with the “generated quantities” section of the syntax

```
1 generated quantities{
2   real slopeG2;
3   slopeG2 = betaHeight + betaHxG2;
4 }
```

The generated quantities block computes values that do not affect the posterior distributions of the parameters—they are computed after the sampling from each iteration

- The values are then added to the Stan object and can be seen in the summary
  - They can also be plotted using the [bayesplot](#) package

```
1 model04b_Samples$summary()
```

```
# A tibble: 9 × 10
  variable      mean  median    sd    mad    q5    q95  rhat ess_bulk
  <chr>      <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>   <dbl>
1 lp__      -150.   -149.   2.05  1.90 -153.  -147.   1.00   1813.
2 beta0        0.214    0.207  0.997  0.993  -1.42   1.86   1.00   4119.
3 betaHeight   12.1    12.0   3.54   3.45   6.33  17.9   1.00  2311.
4 betaGroup2  123.    123.  28.2  28.0  76.3  169.   1.00  3443.
```

# Computing the Slope with Matrices

To put this same method to use with our matrix syntax, we can form a contrast matrix

- Contrasts are linear combinations of parameters
  - You may have used these in R using the [glht](#) package

For us, we form a contrast matrix that is size  $C \times P$  where  $C$  are the number of contrasts

- The entries of this matrix are the values that multiply the coefficients
  - For  $(\beta_1 + \beta_4)$  this would be
    - A one in the corresponding entry for  $\beta_1$
    - A one in the corresponding entry for  $\beta_4$
    - Zeros elsewhere
- $\mathbf{C} = \begin{bmatrix} 0 & 1 & 0 & 0 & 1 & 0 \end{bmatrix}$

The contract matrix then multiplies the coefficients vector to form the values:

$$\mathbf{C}\boldsymbol{\beta}$$

# Contrasts in Stan

We can change our Stan code to import a contrast matrix and use it in generated quantities:

```
1 data {  
2   int<lower=0> N;           // number of observations  
3   int<lower=0> P;           // number of predictors (plus column for intercept)  
4   matrix[N, P] X;          // model.matrix() from R  
5   vector[N] y;             // outcome  
6  
7   vector[P] meanBeta;      // prior mean vector for coefficients  
8   matrix[P, P] covBeta;    // prior covariance matrix for coefficients  
9  
10  real sigmaRate;          // prior rate parameter for residual standard deviation  
11  
12  int<lower=0> nContrasts;  
13  matrix[nContrasts,P] contrast; // contrast matrix for additional effects  
14 }
```

The generated quantities would then become:

```
1 generated quantities {  
2   vector[nContrasts] contrasts;  
3   contrasts = contrastMatrix*beta;  
4 }
```

See example syntax for a full demonstration



# Computing $R^2$

We can use the generated quantities section to build a posterior distribution for  $R^2$

There are several formulas for  $R^2$ , we will use the following:

$$R^2 = 1 - \frac{RSS}{TSS} = 1 - \frac{\sum_{p=1}^N (y_p - \hat{y}_p)^2}{\sum_{p=1}^N (y_p - \bar{y}_p)^2}$$

Where:

- RSS is the regression sum of squares
- TSS is the total sum of squares
- $\hat{y} = \mathbf{X}\beta$
- $\bar{y} = \sum_{p=1}^N \frac{y_p}{N}$

Notice: RSS depends on sampled parameters—so we will use this to build our

posterior distribution for  $R^2$

# Computing $R^2$ in Stan

The generated quantities block can do everything we need to compute  $R^2$

```
1 generated quantities {  
2   vector[nContrasts] heightSlopeG2;  
3   real rss;  
4   real totalrss;  
5  
6   heightSlopeG2 = contrast*beta;  
7  
8   { // anything in these brackets will not appear in summary  
9     vector[N] pred;  
10    pred = X*beta;  
11    rss = dot_self(y-pred); // dot_self is a stan function  
12    totalrss = dot_self(y-mean(y));  
13  }  
14  
15  real R2;  
16  
17  R2 = 1-rss/totalrss;  
18  
19 }
```

See the example syntax for a demonstration

# Wrapping Up

Today we further added to our Bayesian toolset:

- How to make Stan use less syntax using matrices
- How to form posterior distributions for functions of parameters

We will use both of these features in psychometric models

# Up Next

We have one more lecture on linear models that will introduce

- Methods for relative model comparisons
- Methods for checking the absolute fit of a model

Then all things we have discussed to this point will be used in our psychometric models